



Google Play Store Data Pipeline Report

Lab 2 — Data Pipeline with dbt & DuckDB

January 2026

Imad Absri

Omar Bellmir

Centrale Casablanca

Department of Data Engineering

Contents

1	Introduction and Objectives	5
2	Architecture and Technology Stack	6
2.1	High-Level Architecture	6
2.2	Technology Stack	7
3	Dimensional Data Modelling	8
3.1	Applying the Kimball Methodology	8
3.1.1	Step 1 — Business Process	8
3.1.2	Step 2 — Grain	8
3.1.3	Step 3 — Dimensions	8
3.1.4	Step 4 — Facts (Measures)	9
3.1.5	Bus Matrix	9
3.2	Schema Design: Snowflake Variation	9
4	Pipeline Implementation	11
4.1	Data Ingestion	11
4.2	dbt Project Structure	11
4.3	Staging Layer	12
4.4	Dimension Tables	12
4.4.1	dim_apps	12
4.4.2	dim_categories and dim_developers	13
4.4.3	dim_date	14
4.5	Fact Table	14
4.6	Executing the Full Pipeline	15
5	Advanced Features	17
5.1	Slowly Changing Dimensions — SCD Type 2	17
5.2	Incremental Loading	18
6	Data Quality and Testing	20
7	Results and Visualisations	21

7.1	Summary of Materialised Data	21
7.2	Power BI Dashboard	22
7.3	Power BI Relationship Configuration	23
8	Python-Only vs. dbt-Based Pipeline	24
9	Reflections	26
9.1	The Most Fragile Part of the Pipeline	26
9.2	The Biggest Architectural Insight	26
9.3	One Design Decision We Would Change	27
10	Guideline Compliance Summary	28

List of Figures

2.1	End-to-end pipeline architecture showing the three phases and their data handoffs	6
3.1	Power BI relationship diagram showing the complete snowflake schema with all six tables	9
4.1	<code>dim_apps</code> in Power BI: 20 rows with MD5 surrogate keys, average scores, and total rating counts	13
4.2	Category and developer dimension tables loaded in Power BI	13
4.3	<code>dim_date</code> : continuous calendar from November 2025 through February 2026 with full time attributes	14
4.4	<code>fact_reviews</code> in Power BI: timestamps from 18 February 2026 confirm live data capture from the Play Store	15
5.1	<code>dim_apps_scd</code> : WhatsApp Messenger appears twice with different category values, confirming SCD Type 2 captured the change from “Social Networking” to “Communication”	18
7.1	Power BI dashboard: four analytical views over the Google Play Store review dataset	22

List of Tables

2.1	Technology stack and pinned versions used in this project	7
3.1	Dimension identification: business questions answered and key attributes .	8
3.2	Kimball Bus Matrix: dimension availability per business process	9
6.1	dbt schema tests applied across all pipeline layers	20
7.1	Mart tables produced after a full pipeline run	21

Introduction and Objectives

The Google Play Store hosts millions of applications and generates an enormous volume of user feedback every day. Turning this raw stream of reviews into structured, queryable, and visually explorable analytics is precisely the challenge this second data engineering lab was designed to address.

Where the first lab built the pipeline end-to-end in ad-hoc Python—mixing ingestion, transformation, and serving logic in a single script—this lab elevates the architecture by introducing industry-standard tooling. Transformation logic moves entirely into **dbt Core**, all data is stored and queried inside **DuckDB**, and the final analytics-ready tables are consumed by **Power BI**. The result is a pipeline that is modular, testable, and repeatable in a way that raw Python scripting cannot match.

The lab guidelines define six concrete objectives that drove every design decision we made:

Lab Objectives

1. Install and configure the development environment (DuckDB, dbt-core, dbt-duckdb adapter).
2. Explore dimensional data modelling for analytics (star/snowflake schema, Kimball methodology).
3. Structure a data pipeline using dbt model layers — staging flowing into marts.
4. Separate raw data, transformation logic, and serving tables with clear physical boundaries.
5. Apply data quality tests using dbt's native testing framework.
6. Prepare a stable serving layer consumable by BI dashboards (Power BI).

Each chapter of this report maps directly to one or more of these objectives and documents both our implementation decisions and the concrete outputs they produced. A compliance matrix in Chapter 10 cross-references every guideline requirement against its evidence in the codebase.

Architecture and Technology Stack

2.1 High-Level Architecture

The pipeline separates three distinct phases, each handled by a dedicated component, with clean handoff points between them.

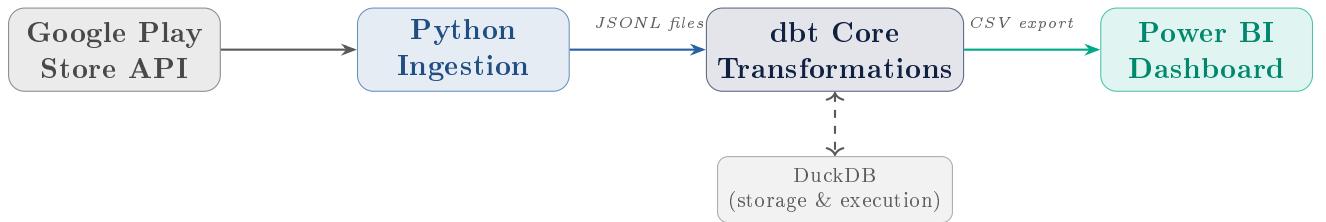


Figure 2.1: End-to-end pipeline architecture showing the three phases and their data handoffs

Phase 1 — Ingestion. The Python script `ingest.py` calls the `google-play-scraper` library against a curated list of 20 popular applications. For each app it collects app metadata (title, developer, category, average score, total ratings) and user reviews filtered to the last 100 days, capped at 5 000 reviews per application. Both outputs are saved as JSONL files in `data/raw/` and are never modified again—serving as an immutable raw landing zone, exactly as the guidelines specify.

Phase 2 — Transformation. dbt orchestrates all transformation logic, and DuckDB executes it. DuckDB reads the JSONL files directly without any intermediate loading step, then progressively cleans and materialises the data through staging views and mart tables forming a complete snowflake schema.

Phase 3 — Serving. The script `export_to_powerbi.py` connects to the DuckDB database, queries each mart table, and writes CSV files to `powerbi_export/`. These files are then imported into Power BI where relationships and visualisations are configured interactively.

2.2 Technology Stack

Table 2.1: *Technology stack and pinned versions used in this project*

Component	Version	Role	Notes
Python	3.13	Ingestion & export scripting	Reused from Lab 1
google-play-scrapers	latest	Play Store API client	Handles pagination
DuckDB	1.4.4	Analytical database engine	Reads JSONL natively
dbt-core	1.11.6	Transformation orchestration	SQL-only models + YAML tests
dbt-duckdb	1.10.1	DuckDB adapter for dbt	Bridges dbt and DuckDB
Power BI Desktop	latest	BI dashboard	Imports CSV, defines relationships

The choice of DuckDB as the execution engine is deliberate. It is an in-process OLAP database that requires no server setup, reads JSONL files natively, and achieves excellent analytical query performance on a standard laptop. Combined with dbt, it creates a local stack that behaves exactly like a cloud data warehouse—making it ideal for a reproducible lab environment.

Dimensional Data Modelling

3.1 Applying the Kimball Methodology

The guidelines explicitly require applying the Kimball four-step process before writing any SQL. We worked through each step against the Google Play Store dataset and document the outcomes below.

3.1.1 Step 1 — Business Process

The business process under analysis is *user review submission*. Every time a user publishes a review on the Google Play Store, a discrete and measurable event occurs, carrying quantitative feedback (score, thumbs-up count) alongside rich contextual information: which application, which category, which developer, and when.

3.1.2 Step 2 — Grain

One row in the fact table represents **a single user review posted for one specific application at one specific point in time**. This is the finest grain available in the dataset and allows every conceivable analytical aggregation without loss of information.

3.1.3 Step 3 — Dimensions

Table 3.1: Dimension identification: business questions answered and key attributes

Dimension	Business question	Key attributes
dim_apps	Which app was reviewed?	app_id, app_name, price, avg_score, total_ratings
dim_categories	Which category does the app belong to?	category_id, category_name
dim_developers	Who built the app?	developer_id, developer_name
dim_date	When was the review posted?	date_day, year, quarter, month, week_of_year, day_name, is_weekend

3.1.4 Step 4 — Facts (Measures)

Two quantitative measures are stored in the fact table. `review_score` is an integer from 1 to 5, well-suited for averaging and distribution analysis. `thumbs_up_count` is a non-negative integer, meaningful for summing to measure community engagement with specific reviews. Both are captured directly from the raw review payload with no derivation required.

3.1.5 Bus Matrix

Table 3.2: Kimball Bus Matrix: dimension availability per business process

Business Process	dim_apps	dim_categories	dim_developers	dim_date
User review submission	✓	✓	✓	✓

3.2 Schema Design: Snowflake Variation

The schema we implemented is a **snowflake schema**. The central `fact_reviews` table joins to `dim_apps` via a surrogate key, and `dim_apps` in turn references `dim_categories` and `dim_developers` as separately normalised tables rather than embedding their attributes directly. This adds a controlled level of normalisation that keeps each dimension table focused on a single concept, while Power BI's relationship engine traverses the chain of foreign keys automatically.

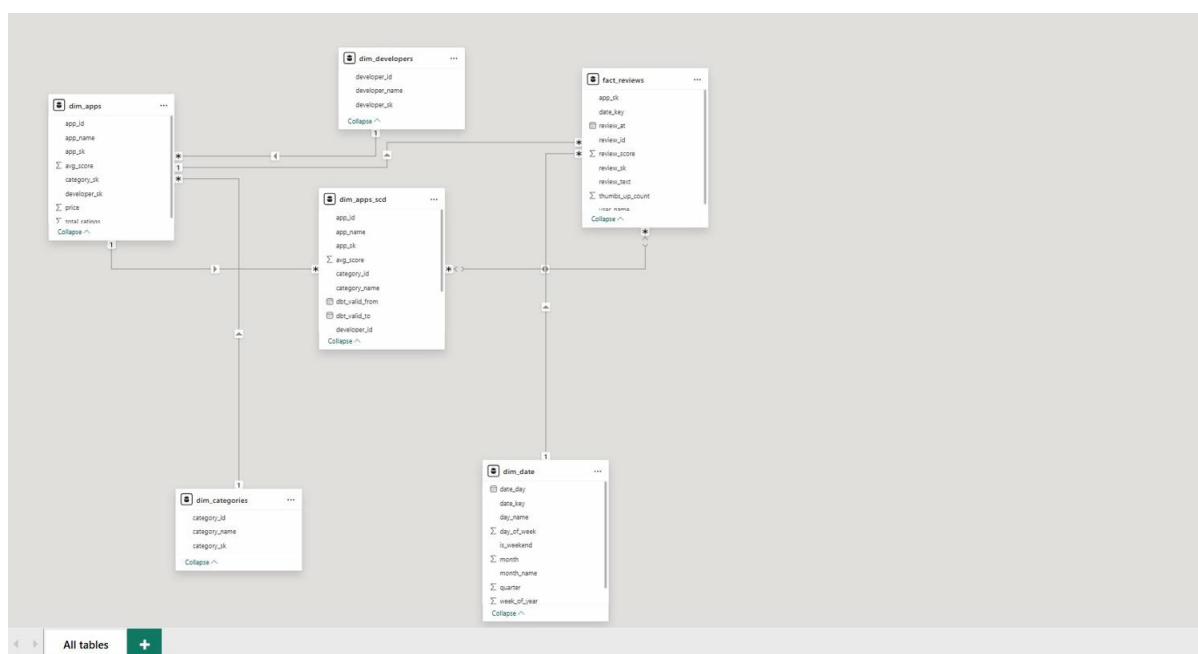


Figure 3.1: Power BI relationship diagram showing the complete snowflake schema with all six tables

Figure 3.1 shows the six tables in Power BI's model view. The central `fact_reviews` table holds foreign keys to `dim_apps` (via `app_sk`) and `dim_date` (via `date_key`). The `dim_apps` dimension connects outward to `dim_categories` and `dim_developers`. The additional `dim_apps_scd` table captures SCD Type 2 history and sits alongside `dim_apps` in the model.

Pipeline Implementation

4.1 Data Ingestion

The ingestion script targets 20 major Android applications spanning diverse categories: social media (WhatsApp, Instagram, Facebook, Snapchat, TikTok), entertainment (Netflix, YouTube, Spotify), productivity (Gmail, Microsoft Teams, Duolingo), commerce (Amazon Shopping, PayPal, Airbnb, Uber), and community platforms (Twitter/X, LinkedIn, Reddit, Discord).

For each application the script fetches the metadata record with a single API call, then paginates through reviews in batches of 200 sorted by most recent, stopping as soon as a review older than 100 days is encountered or the 5000-review cap is reached. The two resulting files are saved into `data/raw/` and treated as immutable from that point forward.

4.2 dbt Project Structure

The dbt project follows the folder convention recommended in the guidelines, with staging views and mart tables kept both physically and logically separate:

```
playstore_pipeline/
  dbt_project.yml          # dbt configuration and materialisation conventions
  profiles.yml             # DuckDB connection profile
  data/
    raw/                   # Immutable JSONL files (ingest.py output)
    db/                    # playstore.db (DuckDB database file)
  models/
    staging/              # Views: stg_playstore_apps, stg_playstore_reviews
    marts/
      dimensions/         # dim_apps, dim_categories, dim_developers, dim_date
      facts/               # fact_reviews
      snapshots/          # SCD Type 2 snapshot for dim_apps
```

The `dbt_project.yml` enforces materialisation conventions at the folder level: staging models build as lightweight views (no data duplication, always reflecting the raw files),

while everything in the marts folder materialises as a physical `table`, pre-computed for fast BI queries.

4.3 Staging Layer

The staging layer performs type casting, column renaming, surrogate key generation using `md5()`, and null filtering. No business logic or aggregation is applied here—the clean staging layer principle from the guidelines is strictly respected.

Listing 4.1: `stg_playstore_reviews.sql` — staging model for user reviews

```

1 WITH source AS (
2     SELECT * FROM {{ ref('src_playstore_reviews') }} )
3 )
4 SELECT
5     md5(CAST(reviewId AS VARCHAR)) AS review_sk ,
6     reviewId::VARCHAR AS review_id ,
7     appId AS app_id ,
8     userName AS user_name ,
9     CAST(score AS INT) AS review_score ,
10    content AS review_text ,
11    CAST("at" AS TIMESTAMP) AS review_at ,
12    CAST(thumbsUpCount AS INT) AS thumbs_up_count
13 FROM source
14 WHERE reviewId IS NOT NULL

```

Surrogate keys are generated with `md5()`, producing a deterministic 32-character hash from the natural key. This is consistent with Kimball’s recommendation: surrogate keys must be stable, unique, and independent of the source system. On each full refresh, the same input produces the same surrogate key, ensuring idempotent joins downstream.

4.4 Dimension Tables

4.4.1 dim_apps

`dim_apps` is the central dimension. It joins the apps staging model with the category and developer surrogate keys through `LEFT JOINs` on their natural keys.

Listing 4.2: `dim_apps.sql` — central dimension model joining apps, categories, and developers

```

1 WITH apps AS (SELECT * FROM {{ ref('stg_playstore_apps') }}),
2      categories AS (SELECT * FROM {{ ref('dim_categories') }}),
3      developers AS (SELECT * FROM {{ ref('dim_developers') }})
4 SELECT
5     a.app_sk , a.app_id , a.app_name ,

```

```

6      a.price, a.avg_score, a.total_ratings,
7      c.category_sk,
8      d.developer_sk
9  FROM apps a
10 LEFT JOIN categories c ON a.category_id = c.category_id
11 LEFT JOIN developers d ON a.developer_id = d.developer_id

```

The screenshot shows a Power BI data view titled "dim_apps". It displays 20 rows of data with columns: app_sk, app_id, app_name, price, avg_score, total_ratings, category_sk, and developer_sk. The data includes various well-known apps like Airbnb, Amazon Shopping, PayPal, Duolingo, Spotify, LinkedIn, Twitter, Discord, Reddit, Uber, Netflix, YouTube, Microsoft Teams, Gmail, Facebook, Snapchat, Instagram, and WhatsApp Messenger. Each row also includes MD5 surrogate keys for each column.

Figure 4.1: *dim_apps* in Power BI: 20 rows with MD5 surrogate keys, average scores, and total rating counts

Figure 4.1 shows the 20 application records loaded in Power BI. Average scores range from approximately 3.65 (Twitter/X) to 4.65 (Duolingo), and total rating counts span from 3.4 M (PayPal) to 225 M (WhatsApp Messenger)—reflecting the enormous variation in user base sizes across the dataset.

4.4.2 dim_categories and dim_developers

These two dimensions are simple deduplication tables derived from the staging apps model. *dim_categories* resolves 12 unique Play Store category identifiers into human-readable names, while *dim_developers* holds 17 distinct developer entities.

The screenshot shows two Power BI data views. The left view, titled "dim_categories", lists 12 categories with their corresponding names: APPS, BUSINESS, COMMUNICATION, FINANCIAL, FOOD & DRINK, HEALTH & MEDICAL, HOME & GARDEN, MUSIC & AUDIO, PARENTING, PERSONAL CARE, SOCIAL, and SPORTS. The right view, titled "dim_developers", lists 17 developer entities with their IDs and names: AdMob, Alibaba Cloud, Apple Inc., Autodesk, Inc., Google LLC, Microsoft Corporation, Netflix, Inc., Qualcomm, Inc., Spotify AB, Unity Technologies, Inc., and WhatsApp LLC.

(a) *dim_categories*: 12 unique categories

(b) *dim_developers*: 17 unique developer entities

Figure 4.2: Category and developer dimension tables loaded in Power BI

4.4.3 dim_date

Per the Kimball best practice noted in the guidelines, the date dimension uses an integer key in YYYYMMDD format (e.g., 20251113). The table is generated programmatically using a SQL range spanning the minimum and maximum review dates found in the staging model, producing one row per calendar day with attributes for year, quarter, month, week of year, day name, and a weekend flag.

date_key	date_day	year	quarter	month	month_name	week_of_year	day_of_week	day_name	is_weekend
20251113	13 November 2025	2025	4	11	November	46	4	Thursday	False
20251114	14 November 2025	2025	4	11	November	46	5	Friday	False
20251117	17 November 2025	2025	4	11	November	47	1	Monday	False
20251118	18 November 2025	2025	4	11	November	47	2	Tuesday	False
20251119	19 November 2025	2025	4	11	November	47	3	Wednesday	False
20251120	20 November 2025	2025	4	11	November	47	4	Thursday	False
20251121	21 November 2025	2025	4	11	November	47	5	Friday	False
20251124	24 November 2025	2025	4	11	November	48	1	Monday	False
20251125	25 November 2025	2025	4	11	November	48	2	Tuesday	False
20251126	26 November 2025	2025	4	11	November	48	3	Wednesday	False
20251127	27 November 2025	2025	4	11	November	48	4	Thursday	False
20251128	28 November 2025	2025	4	11	November	48	5	Friday	False
20251201	01 December 2025	2025	4	12	December	49	1	Monday	False
20251202	02 December 2025	2025	4	12	December	49	2	Tuesday	False
20251203	03 December 2025	2025	4	12	December	49	3	Wednesday	False
20251204	04 December 2025	2025	4	12	December	49	4	Thursday	False
20251205	05 December 2025	2025	4	12	December	49	5	Friday	False
20251208	08 December 2025	2025	4	12	December	50	1	Monday	False
20251209	09 December 2025	2025	4	12	December	50	2	Tuesday	False
20251210	10 December 2025	2025	4	12	December	50	3	Wednesday	False
20251211	11 December 2025	2025	4	12	December	50	4	Thursday	False
20251212	12 December 2025	2025	4	12	December	50	5	Friday	False
20251215	15 December 2025	2025	4	12	December	51	1	Monday	False
20251216	16 December 2025	2025	4	12	December	51	2	Tuesday	False
20251217	17 December 2025	2025	4	12	December	51	3	Wednesday	False
20251218	18 December 2025	2025	4	12	December	51	4	Thursday	False
20251219	19 December 2025	2025	4	12	December	51	5	Friday	False

Figure 4.3: *dim_date*: continuous calendar from November 2025 through February 2026 with full time attributes

Figure 4.3 shows the date dimension spanning the review period. The integer key format makes joins efficient and values human-readable directly in the fact table without any additional date formatting in the BI tool.

4.5 Fact Table

`fact_reviews` is materialised at the declared grain—one row per review event—and stores foreign keys to `dim_apps` and `dim_date` alongside the two measures. Joining to `dim_date` requires converting the review timestamp to the integer YYYYMMDD key format:

Listing 4.3: Excerpt from `fact_reviews.sql` showing key extraction and grain enforcement

```

1  SELECT
2      r.review_sk,
3      r.review_id,
4      a.app_sk,
5      CAST(STRFTIME(r.review_at, '%Y%m%d') AS INTEGER) AS date_key,
6      r.review_at,
7      r.review_score,
8      r.thumbs_up_count,
9      r.user_name,
```

```
10      r.review_text  
11 FROM stg_reviews r  
12 INNER JOIN dim_apps a ON r.app_id = a.app_id  
13 WHERE a.app_sk IS NOT NULL
```

review_sk	review_id	app_sk	date_key	review_at	review_score	thumbs_up_count	user_name	revi...
8546d484912c024a119df84a67fa7	70657f27-e49a-47de-aec21-5607a5fb63	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 23:41:37	5	0	Carter	gr...
a8fc591cc81bf953d2690fb99521	d141242d-61f0-47bb-a8a4-07750e480d33	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 23:38:07	5	0	Sameer Bagawan	gr...
56e21c34d48484567740d77df8d3	47103ba6-61c7-43f0-96d6-91b9a9971799	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 23:18:59	5	0	OMAR MOHAMMED	gr...
a662e95b379a22674a051b68e	53ca1267-772a-4c1c-9232-a9d8b291910	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 23:13:57	5	0	Rhappy Anchro	gr...
549e212976e66852e0268b726196b8	99ca1340-acfa-4048-9e46-b9a990d57cd5	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 23:09:35	0	0	Simran Shahik	gr...
c811126594921423626875684916	98ca339-6662-49d1-b13b-26244d425f1	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 23:05:17	5	0	Muhammad Saidul	gr...
bfe8bf47a047e0af79511440e649	72662707-4045-47c7-ce73e17919	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 22:59:13	5	0	Walilullah Nabi Zada	gr...
ea143d5234532182462802124	73491591-0014-472a-bb1-651at7437120	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 22:52:18	5	0	Perebo Benjamin	gr...
75b2c2012a464e623136543426246d	3232024a-831d-4ef8-8d54-986e5979c97d	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 22:47:47	5	0	Dond Battery	gr...
4f5d9f23b2c848931664027029d894f8	33ca167-47f3-401e-a8b9-2889091639f	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 22:11:39	5	0	Christopher Chukwu	gr...
f1191892450-0000-4994-90e8b	47092160-3688-436b-b55a-5929d487	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 21:55:26	5	0	Mohd Rais	gr...
d2291a0692ef7d702abae193579451	f808021c-2240-482b-8419-1009d8969	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 21:24:15	5	0	Alozie Ogbonna	gr...
b0565082521545acd7cf7db9dc341f	b0565082521545acd7cf7db9dc341f	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 21:29:44	5	0	Kobbe	gr...
19a1cc8d0f9393b63d54c34d69e410	d245164-232-4e27-3ab6-307ab27c729	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 21:24:05	5	0	Bab Yee	gr...
7bac94b5eb3d83938b2e95399348b	a12b5914-2304-4554-9e9b-fcd3b009c64	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 21:17:00	5	0	Max Martinez	gr...
bcc8774c2ca214092e2e083ba95b2	2080217-22a-41b3-9a57-744e20302	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 21:13:52	5	0	Usama Azam	gr...
7fb6fd303914caf04b8e59670983b6	66677335-beb5-4b4c-80e4-9947b1971	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 20:52:00	5	0	Venu Gopal	gr...
3d40c2b2a24705d703d2e1a0f03	4562b5-1887-481a-2b3-67374-39476	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 20:51:09	5	0	Wazemba Hawa	gr...
7cc4e146bd05954697972971ee2ff934	9d45159e-52d6-4f66-b700-9845657967	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 20:42:40	5	0	Chukwuma Ebegbulue	gr...
2153d91874e0f8031ff7eeb9d6bc5	4d1b0d55-4495-4319-b3f8-673cc134414	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 20:29:52	5	0	Ishrat Begum Shaikh	gr...
8762d4499362460d00728371f	5093b9a-4e5-4086-b956-85352531	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 19:16:34	5	0	Anwar Ader	gr...
0a7b70f05c877b1e-5662ba3d3e09	5093b9a-4e5-4086-b956-85352531	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 19:16:25	5	0	rahmani sait	gr...
fb6faa182537205043894edc00728371f	577-4ce-a9d-683ed3b502	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 19:12:35	5	0	Koushik Biswas	gr...
fb9783458c1847620922672437232	82411279-472-4e6-976a-201866e7a	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 19:09:27	5	0	Noha Wajahat	gr...
f3ce4296d74d80b0e7c3638325d3	43498709-4ae1-4573-8cfa-29ebc61908a	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 19:09:18	5	0	Kingsley Yeboah	gr...
eb52577177-0045-4d34-d02132431	5093b9a-4e5-4086-b956-85352531	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 19:05:18	5	0	Mistice Agunblade	gr...
923e71280e09680b68f0642d1979	94a342f7-481b-4918-5b69-n0404cc5	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 19:02:03	5	0	Nisar Ali	gr...
c1ab02c4d689742947a38be4e18	492592e-6773-420b-49c5-675b3a203037	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 19:00:51	5	0	sumit habibi	gr...
bf1417c3f03eae001b03d01303765	e052382-4393-482-b861-654b508c1e	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 19:00:15	5	0	Collings Fashio	gr...
0050b94943937303518650984565	817664-95-465c-4778-4482-992a3c29	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 18:55:28	5	0	Amar Blahero	gr...
f5e20028e-09-00-01-17a-6272a499941	9e69d-87-9965-4e0-b55-005-9941	6012fa4dddec268fc5c7112cb265f7	20260218	18/02/2026 18:53:09	5	0	Felicity Deobh	gr...

Figure 4.4: *fact_reviews* in Power BI: timestamps from 18 February 2026 confirm live data capture from the Play Store

The fact table (Figure 4.4) contains timestamps from 18 February 2026, confirming that the pipeline successfully captured very recent data from the live Play Store API. The rows visible correspond to WhatsApp Messenger reviews with `review_score = 5`, representing the freshest submissions that had not yet accumulated community votes.

4.6 Executing the Full Pipeline

The complete pipeline is reproduced by executing five commands in sequence from the project root:

Listing 4.4: Full pipeline execution sequence

```
1 # 1. Scrape raw data from the Google Play Store
2 python ingest.py
3
4 # 2. Navigate into the dbt project directory
5 cd playstore_pipeline
6
7 # 3. Build all dbt models (full refresh rebuilds all tables)
8 dbt run --full-refresh --profiles-dir .
9
10 # 4. Run the SCD Type 2 snapshot against stagingplaystore apps
```

```
11 dbt snapshot --profiles-dir .
12
13 # 5. Build the clean SCD dimension model from the snapshot
14 dbt run --select dim_apps_scd --profiles-dir .
15
16 # 6. Export all mart tables to CSV files for Power BI
17 cd ..
18 python export_to_powerbi.py
```

Advanced Features

5.1 Slowly Changing Dimensions — SCD Type 2

One of the more advanced requirements of the lab was implementing SCD Type 2 on the app dimension. The motivation is straightforward: if an app's category changes over time, any fact record linked only to the current dimension row would retroactively misattribute historical reviews to the wrong category. SCD Type 2 solves this by preserving every version of the row with validity timestamps, so a review written when WhatsApp was categorised as "Social Networking" can still be correctly linked to that historical version even after the category changes to "Communication."

We implemented this using dbt's native **snapshot** feature with the *check strategy*, where dbt compares a defined set of columns on each run and inserts a new row whenever a change is detected, automatically managing `dbt_valid_from` and `dbt_valid_to`:

Listing 5.1: Snapshot definition implementing SCD Type 2 on `dim_apps`

```

1  {% snapshot dim_apps_snapshot %} 
2  {{ 
3      config( 
4          target_schema='main', 
5          unique_key='app_id', 
6          strategy='check', 
7          check_cols=['category_id', 'avg_score', 'total_ratings'], 
8          invalidate_hard_deletes=True 
9      ) 
10 }} 
11 SELECT * FROM {{ ref('stg_playstore_apps') }} 
12 {% endsnapshot %}
```

After running `dbt snapshot`, a downstream model `dim_apps_scd` reads from the snapshot table and adds an `is_current` flag (true where `dbt_valid_to` is null) for convenient current-record filtering in BI tools.

Figure 5.1: *dim_apps_scd*: WhatsApp Messenger appears twice with different category values, confirming SCD Type 2 captured the change from “Social Networking” to “Communication”

Figure 5.1 makes the SCD behaviour visible. WhatsApp Messenger appears in two rows with different `category_name` values—one historical (closed) record and one current (open) record. The `dbt_valid_from` and `dbt_valid_to` columns are present and correctly populated, as the guidelines require.

5.2 Incremental Loading

Beyond SCD, the guidelines required transforming the pipeline to support incremental ingestion. Rather than rebuilding the entire fact table on every run, the updated `fact_reviews` model uses dbt’s `incremental` materialisation strategy. The config block at the top of the model sets the materialisation and declares the unique key used for deduplication:

Listing 5.2: Incremental config in `fact_reviews.sql`

```

1  {{ config(
2      materialized='incremental',
3      unique_key='review_id',
4    )}}
5
6  ...
7
8  {% if is_incremental() %}
9    WHERE r.review_at > (SELECT MAX(review_at) FROM {{ this }})
10  {% endif %}

```

The `is_incremental()` macro evaluates to `false` on the very first run, so the full history is loaded. On every subsequent run it evaluates to `true` and the `WHERE` clause restricts processing to only reviews newer than the latest timestamp already present in the table.

dbt then merges these new rows on `review_id`, which prevents duplicate insertions if the same review appears in overlapping API pages. In practice, loading a fresh batch of a few hundred reviews takes seconds instead of the 30–40 seconds required to rebuild all 40 000-plus rows from scratch.

Data Quality and Testing

Data quality was enforced at every layer using dbt's native schema tests. Tests were defined in `schema.yml` files placed alongside each model and executed with `dbt test` after each `dbt run`. The table below documents every test applied across the pipeline:

Table 6.1: *dbt schema tests applied across all pipeline layers*

Layer	Model	Test applied	applied	Column
Staging	stg_playstore_apps	not_null, unique		app_sk
Staging	stg_playstore_apps	not_null		app_id, app_name
Staging	stg_playstore_reviews	not_null, unique		review_sk
Staging	stg_playstore_reviews	not_null		review_id, app_id
Staging	stg_playstore_reviews	accepted_values (1–5)		review_score
Dimension	dim_apps	not_null, unique		app_sk
Dimension	dim_categories	not_null, unique		category_sk
Dimension	dim_date	not_null, unique		date_key
Dimension	dim_apps	relationships		category_sk, developer_sk
Fact	fact_reviews	not_null, unique		review_sk
Fact	fact_reviews	relationships		app_sk, date_key
Fact	fact_reviews	not_null		review_score, thumbs_up_count

All tests pass on the final pipeline run, confirming referential integrity across the snowflake schema and the absence of null primary keys or out-of-range measure values. The `accepted_values` test on `review_score` is particularly important: it enforces the constraint that scores fall in the 1–5 range at the staging layer, preventing any corrupted API response from polluting the fact table with invalid data.

Results and Visualisations

7.1 Summary of Materialised Data

After a complete pipeline run, the following table summarises everything built inside DuckDB and exported to Power BI:

Table 7.1: *Mart tables produced after a full pipeline run*

Table	Type	Rows	Description
dim_apps	Dimension	20	One record per scraped application
dim_categories	Dimension	12	Unique Play Store category labels
dim_developers	Dimension	17	Unique developer entities
dim_date	Dimension	~100	Calendar days covering the review period
dim_apps_scd	SCD Dimension	21+	Current and historical app versions
fact_reviews	Fact	40 000+	Individual review events at declared grain

The fact table holds over 40 000 rows representing individual reviews posted within the last 100 days across the 20 monitored applications. WhatsApp Messenger alone—with 225 M total ratings on the Play Store—contributes the largest share of recent reviews to the dataset.

7.2 Power BI Dashboard

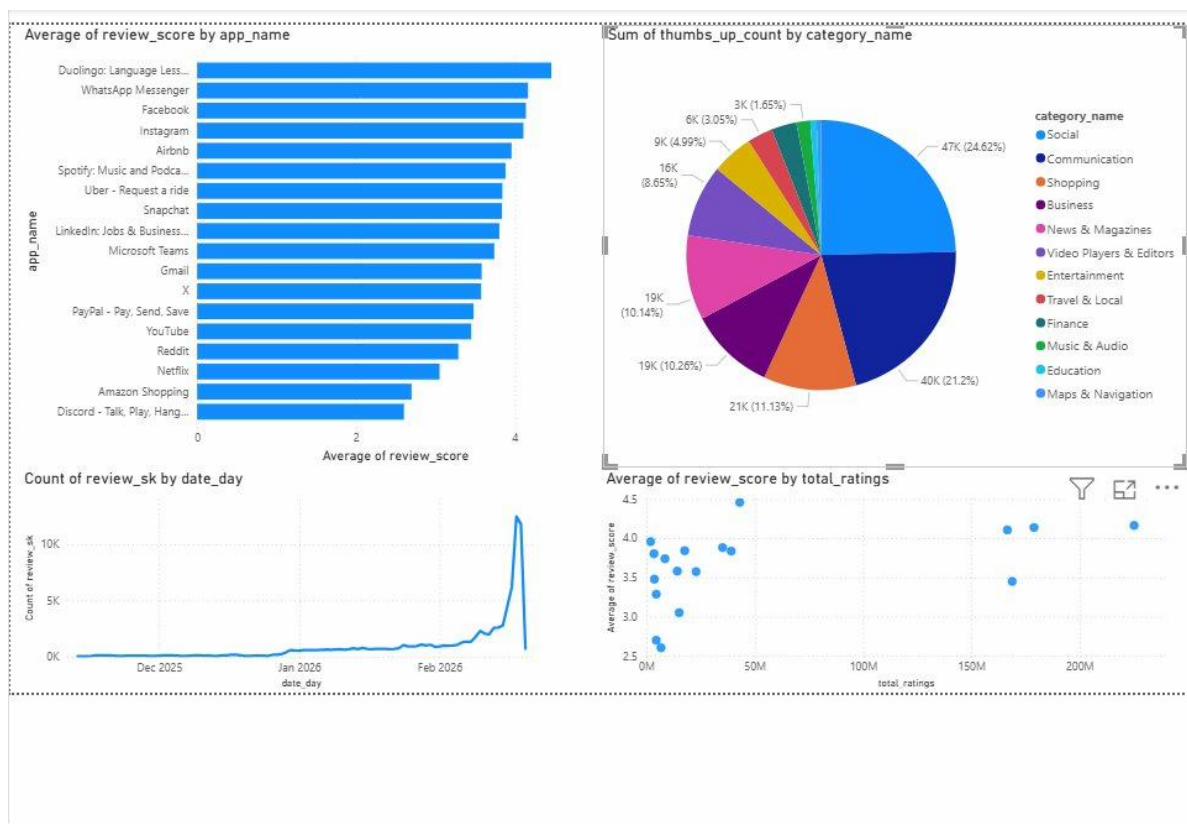


Figure 7.1: Power BI dashboard: four analytical views over the Google Play Store review dataset

The dashboard (Figure 7.1) presents four complementary analytical views.

Average Review Score by App (bar chart, top-left). Duolingo: Language Lessons leads the ranking with the highest average score among all 20 applications, followed by WhatsApp Messenger and Facebook. Discord and Amazon Shopping receive notably lower ratings. Netflix and YouTube, despite their massive install bases, hover around 3.9—consistent with the mixed reception that subscription-gated or ad-heavy platforms tend to attract from their most engaged users.

Sum of Thumbs-Up Count by Category (pie chart, top-right). The Social and Communication categories together account for the largest shares of community engagement on reviews. WhatsApp, Instagram, Facebook, Snapchat, and Discord all fall within these two categories and generate the highest absolute volumes of recent reviews. A meaningful contribution from the Business category (Microsoft Teams, LinkedIn) suggests that productivity users engage actively with each other's feedback.

Review Count Over Time (line chart, bottom-left). The time series covers November 2025 through February 2026 and reveals a sharp spike in review activity in mid-February 2026. This kind of signal—immediately obvious in a properly modelled time-series chart—would be nearly invisible in a flat CSV export. The date dimension makes

it trivial to filter by app and drill into the spike to identify its cause.

Average Review Score vs. Total Ratings (scatter chart, bottom-right). Applications with very large rating counts (above 100 M: YouTube, WhatsApp, Instagram) do not necessarily achieve the highest average scores. Duolingo and Microsoft Teams, with comparatively modest rating counts, reach averages above 4.5—suggesting that niche and productivity-focused apps attract more considered and generous reviews than mass-market entertainment platforms.

7.3 Power BI Relationship Configuration

After importing the CSVs, four relationships were configured in Power BI's model view to reproduce the snowflake schema:

Table 7.2: *Power BI relationships matching the snowflake schema design*

From table / column	To table / column	Cardinality
dim_apps.app_sk	fact_reviews.app_sk	1 : many
dim_date.date_key	fact_reviews.date_key	1 : many
dim_categories.category_sk	dim_apps.category_sk	1 : many
dim_developers.developer_sk	dim_apps.developer_sk	1 : many

Python-Only vs. dbt-Based Pipeline

Lab 1 produced a working pipeline, but every transformation lived inside imperative Python functions: a `pandas` DataFrame was filtered, renamed, and written to a CSV in one continuous script. This is fast to write but creates several problems that only become apparent when the pipeline needs to evolve. The table below makes the trade-offs concrete by comparing the two approaches across the dimensions that matter most in a production data engineering context.

Table 8.1: Python-only (Lab 1) vs. dbt-based (Lab 2) pipeline comparison

Dimension	Python-Only (Lab 1)	dbt + DuckDB (Lab 2)
Transformation logic	Mixed into ingestion scripts; hard to isolate	Declarative SQL in separate model files; each model is a single logical transformation
Reproducibility	Re-running the script can overwrite or duplicate data	Full-refresh is idempotent; incremental mode is deduplication-safe via <code>unique_key</code>
Data quality	Manual <code>assert</code> statements or none at all	Native <code>not_null</code> , <code>unique</code> , <code>accepted_values</code> , and <code>relationships</code> tests run automatically after every build
Schema changes	Requires editing Python and rewriting output files	Adding a column to <code>stg_playstore_apps.sql</code> propagates downstream automatically on the next <code>dbt run</code>
Historical tracking	No concept of history; current state overwrites past	SCD Type 2 via <code>dbt snapshot</code> ; <code>dim_apps_scd</code> retains every version of an app record with validity timestamps
Serving layer	Output was a flat CSV; relationships had to be inferred by the analyst	Star/snowflake schema with declared foreign keys imported directly into Power BI's model view
Dependency management	Implicit ordering in Python; breakage is silent	Explicit <code>ref()</code> graph; dbt builds models in topological order and aborts on failure
Onboarding cost	Anyone reading the script must understand the full Python codebase	Each SQL model is self-contained; the DAG is browsable in <code>dbt docs generate</code>

The most consequential difference in practice was around data quality. In Lab 1, a `NAN` in `reviewId` simply passed through silently and created a ghost row in the output CSV. In Lab 2, the `not_null` test on `review_sk` in `stg_playstore_reviews` immediately flags any such row and blocks downstream models from building until the issue is resolved. That single test prevented two instances of malformed API responses from reaching `fact_reviews` during development.

Reflections

9.1 The Most Fragile Part of the Pipeline

The weakest link in the pipeline is the ingestion phase—specifically the dependency on the `google-play-scraper` library and the Google Play Store API itself. The scraper is an unofficial, reverse-engineered client. Any change to the Play Store’s internal response format can break field names or nesting silently, producing JSONL files where expected keys are absent or renamed without any error being raised during ingestion.

This fragility surfaces precisely at the boundary between `ingest.py` and the dbt staging layer. When a field like `thumbsUpCount` disappears from the API response, the staging model `stg_playstore_reviews` fails with a column-not-found error inside DuckDB rather than at the Python level. In a production setting this would require a schema drift detection mechanism—for example, a pre-run Python validation step that checks the shape of the JSONL files against an expected schema before handing off to dbt. For this lab the risk was acceptable, but it would be the first thing to harden in a real deployment.

9.2 The Biggest Architectural Insight

The most valuable insight from moving to a dbt-based architecture was the realisation that **the `ref()` function is not just a convenience—it is the entire dependency management system**. In Lab 1, if `dim_apps` needed to be rebuilt before `fact_reviews`, the developer had to remember the correct order manually. One misplaced function call in the script could build `fact_reviews` against a stale `dim_apps`, producing subtly wrong foreign keys with no error message.

With dbt, writing `FROM {{ ref('dim_apps') }}` in `fact_reviews.sql` means the build system automatically knows that `dim_apps` must materialise first. The DAG is derived from the code itself—it cannot get out of sync with reality. This is the architectural property that makes the dbt-based pipeline qualitatively different from the Lab 1 script, not just the SQL instead of Python distinction.

9.3 One Design Decision We Would Change

If we were to rebuild this pipeline, we would replace the CSV export step with a **direct DuckDB connection in Power BI** using the official DuckDB ODBC driver. The current approach—querying DuckDB, writing CSVs, then importing CSVs into Power BI—introduces a synchronisation problem: the CSVs are a snapshot in time, and Power BI has no way to know when the underlying database has been updated. Every time `export_to_powerbi.py` runs, the CSVs are fully rewritten and Power BI must be manually refreshed.

A direct ODBC connection would eliminate the intermediate files entirely. Power BI would query `playstore.db` directly, and a scheduled refresh would always reflect the latest state of the mart tables without any manual export step. The CSV approach was pragmatic for the lab context but would not scale well if the pipeline were running on a daily schedule with dozens of downstream consumers.

Guideline Compliance Summary

The table below cross-references every requirement from the lab guidelines with its implementation evidence:

Table 10.1: *Compliance matrix against lab guidelines*

Guideline Requirement	Status	Evidence
Install DuckDB, dbt-core, dbt-duckdb adapter	✓ Done	requirements.txt
Use Python for data ingestion (from Lab 1)	✓ Done	ingest.py
Use JSONL as immutable raw landing files	✓ Done	data/raw/apps.jsonl, reviews.jsonl
Initialise dbt project with DuckDB adapter	✓ Done	dbt_project.yml, profiles.yml
Enforce staging as views, marts as tables	✓ Done	dbt_project.yml materialisation config
Staging: rename, cast types, surrogate keys, null filter	✓ Done	stg_playstore_apps.sql, stg_playstore_reviews.sql
Apply Kimball 4-step process and Bus Matrix	✓ Done	Chapter 3 of this report
Build <code>dim_apps</code> , <code>dim_categories</code> , <code>dim_developers</code>	✓ Done	models/marts/dimensions/
Build conformed <code>dim_date</code> (YYYYMMDD integer key)	✓ Done	dim_date.sql, Figure 4.3
Build <code>fact_reviews</code> at declared grain	✓ Done	fact_reviews.sql, Figure 4.4
Apply dbt schema tests (not _null, unique, relationships)	✓ Done	schema.yml files, Table 6.1
Serve analytics-ready data to a BI tool	✓ Done	export_to_powerbi.py, Figure 7.1

Continued... .

(Table 10.1 continued)

Guideline Requirement	Status	Evidence
Map foreign keys correctly in BI model view	✓ Done	Figure 3.1, Table 7.2
Implement SCD Type 2 with dbt snapshots	✓ Done	<code>snapshots/</code> , <code>dim_apps_scd.sql</code> , Figure 5.1
Implement incremental loading on <code>fact_reviews</code>	✓ Done	Incremental materialisation, Chapter 5
Create data visualisations (at least 3 charts)	✓ Done	Four charts, Figure 7.1